

HIVE PROJECT 1

Q1) Create a schema based on the given dataset

Q2) Dump the data inside the hdfs in the given schema location.

Q1,2) Solution

Make directory at hdfs location.

```
[cloudera@quickstart ~]$ hdfs dfs -mkdir Agent
```

Put files (AgentLoggingReport.csv Agent, AgentPerformance.csv Agent) into Agent directory in hdfs location.

```
[cloudera@quickstart ~]$ hdfs dfs -put /tmp/Agent_data/AgentLoggingReport.csv Agent/
```

```
[cloudera@quickstart ~]$ hdfs dfs -put /tmp/Agent_data/AgentPerformance.csv Agent/
```

Check the files inside hdfs location.

```
[cloudera@quickstart ~]$ hdfs dfs -ls Agent
```

Found 2 items

```
-rw-r--r-- 1 cloudera cloudera 55351 2022-11-01 08:58 Agent/AgentLoggingReport.csv
```

```
-rw-r--r-- 1 cloudera cloudera 109853 2022-11-01 08:59 Agent/AgentPerformance.csv
```

Go to hive and use agent database.

```
hive> use agent;
```

OK

Time taken: 1.25 seconds

Create tables:

```
hive> create table agent_logging
```

```
> (
```

```
> s_no int,
```

```
> agent string,
```

```
> date date,
```

```
> login_time string,
```

```
> logout_time string,
```

```
> duration string
```

```
> )
```

```
> row format delimited
```

```
> fields terminated by ','
```

```
> tblproperties("skip.header.line.count"="1");
```

OK

Time taken: 0.662 seconds

hive> create table agent_performance

```
> (  
> s_no int,  
> date date,  
> agent_name string,  
> total_chats int,  
> average_response_time string,  
> average_resolution_time string,  
> average_rating float,  
> total_feedback int  
> )  
  
> row format delimited  
  
> fields terminated by ','  
  
> tblproperties("skip.header.line.count"="1");
```

OK

Time taken: 0.218 seconds

Load data inside above created tables.

hive> load data inpath 'Agent/AgentLoggingReport.csv' into table agent_logging;

Loading data to table agent.agent_logging

Table agent.agent_logging stats: [numFiles=1, totalSize=56353]

OK

Time taken: 1.581 seconds

hive> load data inpath 'Agent/AgentPerformance.csv' into table agent_performance;

Loading data to table agent.agent_performance

Table agent.agent_performance stats: [numFiles=1, totalSize=116159]

OK

Time taken: 0.837 seconds

Fetch some records from tables.

```
hive> select * from agent_logging limit 3;
```

OK

1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10
2	Khushboo Priya	2022-07-30	15:06:59	15:07:16	00:00:17
3	Nandani Gupta	2022-07-30	15:04:24	17:31:07	02:26:42

Time taken: 0.971 seconds, Fetched: 3 row(s)

```
hive> select * from agent_performance limit 3;
```

OK

1	2022-07-30	Purna Singh	11	00:00:38	00:04:20	4.11	9
2	2022-07-30	Nandani Gupta	11	00:01:15	00:28:25	3.14	7
3	2022-07-30	Ameya Jain	14	00:00:30	00:11:36	4.55	11

Time taken: 0.118 seconds, Fetched: 3 row(s)

Q3) List of all agents' names.

Solution:

```
hive> select distinct agent from agent_logging;
```

Query ID = cloudera_20221101210404_30dc2792-6a5d-46e6-95d7-61106454adf0

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0002/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:05:08,000 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:05:23,698 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.81 sec

2022-11-01 21:05:37,446 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.11 sec

MapReduce Total cumulative CPU time: 6 seconds 110 msec

Ended Job = job_1667316033891_0002

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.11 sec HDFS Read: 64055 HDFS Write: 638 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 110 msec

OK

Aditya Shinde

Aditya _iot

Amersh

Ameya Jain

Ankitjha

Anurag Tiwari

Aravind

Ayushi Mishra

Bharath

Boktiar Ahmed Bappy

Chaitra K Hiremath

Deepranjan Gupta

Dibyanshu

Harikrishnan Shaji

Hrisikesh Neogi

Hyder Abbas

Ineuron Intelligence

Ishawant Kumar

Jawala Prakash

Jaydeep Dixit

Khushboo Priya

Madhulika G

Mahesh Sarade

Maitry

Manjunatha A

Mithun S

Mukesh

Muskan Garg

Nandani Gupta

Nishtha Jain

Nitin M

Prabir Kumar Satapathy

Prateek _iot

Perna Singh

Rishav Dash

Saikumarreddy N

Sanjeev Kumar

Saurabh Shukla

Shiva Srivastava

Shivan K

Shivananda Sonwane

Shubham Sharma

Sowmiya Sivakumar

Sudhanshu Kumar

Suraj S Bilgi

Swati

Tarun

Wasim

Zeeshan

Time taken: 60.25 seconds, Fetched: 49 row(s)

hive> select distinct agent_name from agent_performance;

Query ID = cloudera_20221101210606_286d5cfa-e577-47db-9828-045875c5505f

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0003/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0003

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:06:19,465 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:06:32,624 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.81 sec

2022-11-01 21:06:48,642 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.63 sec

MapReduce Total cumulative CPU time: 6 seconds 630 msec

Ended Job = job_1667316033891_0003

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.63 sec HDFS Read: 124583 HDFS Write: 841 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 630 msec

OK

Abhishek
Aditya
Aditya Shinde
Aditya_iot
Amersh
Ameya Jain
Anirudh
Ankit Sharma
Ankitjha
Anurag Tiwari
Aravind
Ashad Nasim
Ashish
Ayushi Mishra
Bharath
Boktiar Ahmed Bappy
Chaitra K Hiremath
Deepranjan Gupta
Dibyanshu
Harikrishnan Shaji
Hitesh Choudhary
Hrisikesh Neogi
Hyder Abbas
Ineuron Intelligence
Ishawant Kumar
Jawala Prakash
Jayant Kumar
Jaydeep Dixit
Khushboo Priya
Madhulika G
Mahak
Mahesh Sarade
Maitry
Maneesh
Manjunatha A
Mithun S
Mukesh

Mukesh Rao
Muskan Garg
Nandani Gupta
Nishtha Jain
Nitin M
Prabir Kumar Satapathy
Prateek _iot
Perna Singh
Rishav Dash
Rohan
Saif Khan
Saikumarreddy N
Samprit
Sandipan Saha
Sanjeev Kumar
Sanjeevan
Saurabh Shukla
Shiva Srivastava
Shivan K
Shivan_S
Shivananda Sonwane
Shubham Sharma
Sowmiya Sivakumar
Spuri
Sudhanshu Kumar
Suraj S Bilgi
Swati
Tarun
Uday Mishra
Vasanth P
Vivek
Wasim
Zeeshan

Time taken: 46.158 seconds, Fetched: 71 row(s)

Q4) Find out agent average rating.

Solution:

```
hive> set hive.cli.print.header = true;
```

```
hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance
group by agent_name;
```

Query ID = cloudera_20221101211111_d3477195-298c-4632-97c5-05e0bc981a44

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0006/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0006

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:11:37,012 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:11:44,688 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.33 sec

2022-11-01 21:11:57,761 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.79 sec

MapReduce Total cumulative CPU time: 4 seconds 790 msec

Ended Job = job_1667316033891_0006

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.79 sec HDFS Read: 125817 HDFS Write: 1881 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 790 msec

OK

agent average_rating

Abhishek 0.0

Aditya 0.0

Aditya Shinde 1.8003333409627278

Aditya_iot 2.3453333377838135

Amersh 0.0

Ameya Jain 2.21966667175293

Anirudh 0.6449999968210857

Ankit Sharma 0.0

Ankitjha 0.26666666666666666

Anurag Tiwari 0.18333333333333332

Aravind 2.1813333511352537

Ashad Nasim 0.16666666666666666

Ashish 0.0

Ayushi Mishra 3.481999969482422

Bharath 2.9836666584014893

Boktiar Ahmed Bappy 3.567999982833862

Chaitra K Hiremath 0.8646666606267294

Deepranjan Gupta 2.886666695276896

Dibyanshu 0.0

Harikrishnan Shaji 2.6396666526794434

Hitesh Choudhary 0.0

Hrisikesh Neogi 3.1363333304723104

Hyder Abbas 0.0

Ineuron Intelligence 0.0

Ishawant Kumar 3.543333347638448

Jawala Prakash 3.472000018755595

Jayant Kumar 1.068666664759318

Jaydeep Dixit 3.1670000314712525

Khushboo Priya 3.703666663169861

Madhulika G 3.4986666520436605

Mahak 0.1

Mahesh Sarade 2.4003333330154417

Maitry 2.9270000139872234

Maneesh 0.16666666666666666

Manjunatha A 3.5946666876475017

Mithun S 2.359000023206075

Mukesh 0.3096666653951009

Mukesh Rao 0.25566666523615517

Muskan Garg 0.712333329518636

Nandani Gupta 2.9236666679382326

Nishtha Jain 3.282333334287008

Nitin M 0.0

Prabir Kumar Satapathy 2.5103333314259846

Prateek _iot 2.4383333206176756

Prerna Singh 3.2326666434605915

Rishav Dash 1.426833335585734

Rohan 0.0

Saif Khan 0.0

Saikumarreddy N 1.9803333441416422
Samprit 0.0
Sandipan Saha 0.4289999961853027
Sanjeev Kumar 3.3830000241597493
Sanjeevan 0.0
Saurabh Shukla 0.5556666692097981
Shiva Srivastava 0.9446666717529297
Shivan K 2.841333341598511
Shivan_S 0.14166666666666666
Shivananda Sonwane 4.232666659355163
Shubham Sharma 3.2253333568572997
Sowmiya Sivakumar 1.2599999984105428
Spuri 0.0
Sudhanshu Kumar 0.3333333333333333
Suraj S Bilgi 0.31200000445048015
Swati 2.4236666917800904
Tarun 0.05
Uday Mishra 0.0
Vasanth P 0.0
Vivek 0.5006666660308838
Wasim 2.400000015894572
Zeeshan 2.286999988555908
Time taken: 33.155 seconds, Fetched: 71 row(s)

Q5) Total working days for each agent

Solution:

hive> select agent_name as agent, count(distinct date) as number_of_working_days from agent_performance group by agent_name;

Query ID = cloudera_20221101211717_617f0d44-ffd7-46f9-b0a1-e8b032095df8

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0008/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0008

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:17:53,354 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:18:01,959 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.62 sec

2022-11-01 21:18:13,514 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.15 sec

MapReduce Total cumulative CPU time: 5 seconds 150 msec

Ended Job = job_1667316033891_0008

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.15 sec HDFS Read: 125596 HDFS Write: 1053 SUCCESS

Total MapReduce CPU Time Spent: 5 seconds 150 msec

OK

agent number_of_working_days

0

Abhishek 30

Aditya 30

Aditya Shinde 30

Aditya_iot 30

Amersh 30

Ameya Jain 30

Anirudh 30

Ankit Sharma 30

Ankitjha 30

Anurag Tiwari 30

Aravind 30

Ashad Nasim 30

Ashish 30

Ayushi Mishra 30

Bharath 30

Boktiar Ahmed Bappy 30

Chaitra K Hiremath 30

Deepranjan Gupta 30

Dibyanshu 30

Harikrishnan Shaji 30

Hitesh Choudhary 30

Hrisikesh Neogi 30

Hyder Abbas 30

Ineuron Intelligence 30

Ishawant Kumar 30

Jawala Prakash 30

Jayant Kumar 30

Jaydeep Dixit 30

Khushboo Priya 30

Madhulika G 30

Mahak 30

Mahesh Sarade 30

Maitry 30

Maneesh 30

Manjunatha A 30

Mithun S 30

Mukesh 30

Mukesh Rao 30

Muskan Garg 30

Nandani Gupta 30

Nishtha Jain 30

Nitin M 30

Prabir Kumar Satapathy 30

Prateek _iot 30

Perna Singh 30

Rishav Dash 30

Rohan 30

Saif Khan 30

Saikumarreddy N 30

Samprit 30

Sandipan Saha 30

Sanjeev Kumar 30

Sanjeevan 30

Saurabh Shukla 30

Shiva Srivastava 30

Shivan K 30

Shivan_S 30

Shivananda Sonwane 30

Shubham Sharma 30

Sowmiya Sivakumar 30

Spuri 30

Sudhanshu Kumar 30

Suraj S Bilgi 30

Swati 30

Tarun 30

Uday Mishra 30

Vasanth P 30

Vivek 30

Wasim 30

Zeeshan 30

Time taken: 31.571 seconds, Fetched: 71 row(s)

Q6) Total query that each agent has taken

Solution:

hive> select agent_name as agent, sum(total_chats) as queries_taken from agent_performance group by agent_name;

Query ID = cloudera_20221101214848_5e4743f7-bfa2-437e-a8af-18d16fa84039

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0009, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0009/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0009

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:48:24,468 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:48:33,352 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.17 sec

2022-11-01 21:48:44,544 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.69 sec

MapReduce Total cumulative CPU time: 4 seconds 690 msec

Ended Job = job_1667316033891_0009

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.69 sec HDFS Read: 125339 HDFS Write: 1065 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 690 msec

OK

agent queries_taken

Abhishek 0

Aditya 0

Aditya Shinde 277

Aditya_iot 231

Amersh 0

Ameya Jain 322

Anirudh 81

Ankit Sharma 0

Ankitjha 5

Anurag Tiwari 4

Aravind 366

Ashad Nasim 18

Ashish 0

Ayushi Mishra 514

Bharath 369

Boktiar Ahmed Bappy 452

Chaitra K Hiremath 64

Deepranjan Gupta 493

Dibyanshu 1

Harikrishnan Shaji 381

Hitesh Choudhary 1

Hrisikesh Neogi 578

Hyder Abbas 0

Ineuron Intelligence 0

Ishawant Kumar 338

Jawala Prakash 439

Jayant Kumar 127

Jaydeep Dixit 512

Khushboo Priya 446

Madhulika G 469

Mahak 7

Mahesh Sarade 364

Maitry 542

Maneesh 4

Manjunatha A 413

Mithun S 503

Mukesh 19

Mukesh Rao	5
Muskan Garg	56
Nandani Gupta	560
Nishtha Jain	373
Nitin M	0
Prabir Kumar Satapathy	299
Prateek _iot	190
Perna Singh	401
Rishav Dash	409
Rohan	0
Saif Khan	0
Saikumarreddy N	364
Samprit	1
Sandipan Saha	30
Sanjeev Kumar	507
Sanjeevan	0
Saurabh Shukla	16
Shiva Srivastava	53
Shivan K	357
Shivan_S	7
Shivananda Sonwane	441
Shubham Sharma	510
Sowmiya Sivakumar	206
Spuri	0
Sudhanshu Kumar	2
Suraj S Bilgi	28
Swati	524
Tarun	22
Uday Mishra	0
Vasanth P	0
Vivek	44
Wasim	433
Zeeshan	542

Time taken: 30.731 seconds, Fetched: 71 row(s)

Q7) Total Feedback that each agent has received

Solution:

hive> select agent_name as agent, sum(total_feedback) as feedbacks_received from agent_performance group by agent_name;

Query ID = cloudera_20221101215656_81013437-bff7-447c-b446-3709c703263a

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0012, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0012/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0012

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 21:56:17,193 Stage-1 map = 0%, reduce = 0%

2022-11-01 21:56:29,417 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.06 sec

2022-11-01 21:56:42,782 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.62 sec

MapReduce Total cumulative CPU time: 6 seconds 620 msec

Ended Job = job_1667316033891_0012

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.62 sec HDFS Read: 125346 HDFS Write: 1061 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 620 msec

OK

agent feedbacks_received

Abhishek 0

Aditya 0

Aditya Shinde 153

Aditya_iot 131

Amersh 0

Ameya Jain 228

Anirudh 39

Ankit Sharma 0

Ankitjha 3

Anurag Tiwari 3

Aravind 233

Ashad Nasim 9

Ashish 0

Ayushi Mishra 329

Bharath 247

Boktiar Ahmed Bappy 311

Chaitra K Hiremath 37

Deepranjan Gupta 312

Dibyanshu 0

Harikrishnan Shaji 231

Hitesh Choudhary 0

Hrisikesh Neogi 367

Hyder Abbas 0

Ineuron Intelligence 0

Ishawant Kumar 202

Jawala Prakash 250

Jayant Kumar 70

Jaydeep Dixit 305

Khushboo Priya 289

Madhulika G 281

Mahak 5

Mahesh Sarade 216

Maitry 347

Maneesh 3

Manjunatha A 254

Mithun S 364

Mukesh 17

Mukesh Rao 5

Muskan Garg 37

Nandani Gupta 308

Nishtha Jain 257

Nitin M 0

Prabir Kumar Satapathy 222

Prateek _jot 107

Perna Singh 235

Rishav Dash 264

Rohan 0

Saif Khan 0

Saikumarreddy N 290

Samprit 0

Sandipan Saha 18
Sanjeev Kumar 311
Sanjeevan 0
Saurabh Shukla 8
Shiva Srivastava 46
Shivan K 243
Shivan_S 4
Shivananda Sonwane 263
Shubham Sharma 300
Sowmiya Sivakumar 141
Spuri 0
Sudhanshu Kumar 2
Suraj S Bilgi 15
Swati 302
Tarun 6
Uday Mishra 0
Vasanth P 0
Vivek 20
Wasim 284
Zeeshan 335

Time taken: 42.047 seconds, Fetched: 71 row(s)

Q8) Agent name who have average rating between 3.5 to 4

Solution:

hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance group by agent_name having average_rating between 3.5 and 4;

Query ID = cloudera_20221101220202_7dfde0dd-a9b0-4135-ab59-f4fa72fc5b7d

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0013, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0013/

```

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0013

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 22:02:38,620 Stage-1 map = 0%, reduce = 0%

2022-11-01 22:02:48,489 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.04 sec

2022-11-01 22:02:59,284 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.85 sec

MapReduce Total cumulative CPU time: 4 seconds 850 msec

Ended Job = job_1667316033891_0013

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.85 sec HDFS Read: 126449 HDFS Write: 136 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 850 msec

OK

agent average_rating

Boktiar Ahmed Bappy 3.567999982833862

Ishawant Kumar 3.543333347638448

Khushboo Priya 3.703666663169861

Manjunatha A 3.5946666876475017

Time taken: 33.547 seconds, Fetched: 4 row(s)

```

Q9) Agent name who have rating less than 3.5

Solution:

```
hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance
group by agent_name having average_rating < 3.5;
```

```
Query ID = cloudera_20221101220404_673bfe50-feae-4738-b904-6f241f2aa822
```

```
Total jobs = 1
```

```
Launching Job 1 out of 1
```

```
Number of reduce tasks not specified. Estimated from input data size: 1
```

```
In order to change the average load for a reducer (in bytes):
```

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

```
In order to limit the maximum number of reducers:
```

```
set hive.exec.reducers.max=<number>
```

```
In order to set a constant number of reducers:
```

```
set mapreduce.job.reduces=<number>
```

```
Starting Job = job_1667316033891_0014, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0014/
```

```
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0014
```

```
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
```

```
2022-11-01 22:04:34,295 Stage-1 map = 0%, reduce = 0%
```

```
2022-11-01 22:04:46,500 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.71 sec
```

2022-11-01 22:05:03,498 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.32 sec

MapReduce Total cumulative CPU time: 7 seconds 320 msec

Ended Job = job_1667316033891_0014

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.32 sec HDFS Read: 126252 HDFS Write: 1704 SUCCESS

Total MapReduce CPU Time Spent: 7 seconds 320 msec

OK

agent average_rating

Abhishek 0.0

Aditya 0.0

Aditya Shinde 1.8003333409627278

Aditya_jot 2.3453333377838135

Amersh 0.0

Ameya Jain 2.21966667175293

Anirudh 0.6449999968210857

Ankit Sharma 0.0

Ankitjha 0.26666666666666666

Anurag Tiwari 0.18333333333333332

Aravind 2.1813333511352537

Ashad Nasim 0.16666666666666666

Ashish 0.0

Ayushi Mishra 3.481999969482422

Bharath 2.9836666584014893

Chaitra K Hiremath 0.8646666606267294

Deepranjan Gupta 2.886666695276896

Dibyanshu 0.0

Harikrishnan Shaji 2.6396666526794434

Hitesh Choudhary 0.0

Hrisikesh Neogi 3.1363333304723104

Hyder Abbas 0.0

Ineuron Intelligence 0.0

Jawala Prakash 3.472000018755595

Jayant Kumar 1.068666664759318

Jaydeep Dixit 3.1670000314712525

Madhulika G 3.4986666520436605

Mahak 0.1

Mahesh Sarade 2.4003333330154417

Maitry 2.9270000139872234

Maneesh 0.16666666666666666

Mithun S 2.359000023206075

Mukesh 0.3096666653951009

Mukesh Rao 0.25566666523615517

Muskan Garg 0.712333329518636

Nandani Gupta 2.9236666679382326

Nishtha Jain 3.282333334287008

Nitin M 0.0

Prabir Kumar Satapathy 2.5103333314259846

Prateek_jot 2.4383333206176756

Perna Singh 3.2326666434605915

Rishav Dash 1.426833335585734

Rohan 0.0

Saif Khan 0.0

Saikumarreddy N 1.9803333441416422

Samprit 0.0

Sandipan Saha 0.4289999961853027

Sanjeev Kumar 3.3830000241597493

Sanjeevan 0.0

Saurabh Shukla 0.5556666692097981

Shiva Srivastava 0.9446666717529297

Shivan K 2.841333341598511

Shivan_S 0.14166666666666666

Shubham Sharma 3.2253333568572997

Sowmiya Sivakumar 1.2599999984105428

Spuri 0.0

Sudhanshu Kumar 0.3333333333333333

Suraj S Bilgi 0.31200000445048015

Swati 2.4236666917800904

Tarun 0.05

Uday Mishra 0.0

Vasanth P 0.0

Vivek 0.500666660308838

Wasim 2.400000015894572

Zeeshan 2.286999988555908

Time taken: 45.33 seconds, Fetched: 65 row(s)

Q10) Agent name who have rating more than 4.5

Solution:

```
hive> select agent_name as agent, avg(average_rating) as average_rating from agent_performance  
group by agent_name having average_rating > 4.5;
```

Query ID = cloudera_20221101220505_cc78d0b4-99be-4cd0-8227-fe4d79e36f40

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0015, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0015/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0015

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 22:05:30,465 Stage-1 map = 0%, reduce = 0%

2022-11-01 22:05:43,975 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.94 sec

2022-11-01 22:05:59,236 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.62 sec

MapReduce Total cumulative CPU time: 6 seconds 620 msec

Ended Job = job_1667316033891_0015

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.62 sec HDFS Read: 126251 HDFS Write: 0 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 620 msec

OK

agent average_rating

Time taken: 44.717 seconds

Q11) How many feedback agents have received more than 4.5 average

Solution:

```
hive> select count(*) from(select agent_name as agent, avg(total_feedback) as average_feedback  
from agent_performance group by agent_name having average_feedback > 4.5);
```

Query ID = cloudera_20221101221919_c687ebc1-ccb0-4d82-bd3b-2bb333af0f74

Total jobs = 2

Launching Job 1 out of 2

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0019, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0019/

Kill Command = `/usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0019`

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 22:20:02,892 Stage-1 map = 0%, reduce = 0%

2022-11-01 22:20:14,116 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.71 sec

2022-11-01 22:20:27,703 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.64 sec

MapReduce Total cumulative CPU time: 6 seconds 640 msec

Ended Job = job_1667316033891_0019

Launching Job 2 out of 2

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0020, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0020/

Kill Command = `/usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0020`

Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1

2022-11-01 22:20:42,088 Stage-2 map = 0%, reduce = 0%

2022-11-01 22:20:52,198 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.39 sec

2022-11-01 22:21:04,433 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 5.45 sec

MapReduce Total cumulative CPU time: 5 seconds 450 msec

Ended Job = job_1667316033891_0020

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.64 sec HDFS Read: 125695 HDFS Write: 114 SUCCESS

Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 5.45 sec HDFS Read: 4550 HDFS Write: 3 SUCCESS

Total MapReduce CPU Time Spent: 12 seconds 90 msec

OK

_c0

31

Time taken: 77.471 seconds, Fetched: 1 row(s)

Q12) average weekly response time for each agent

Solution:

```
hive> select agent, avg(weekly_response_time_in_sec) as avg_weekly_response_time_in_sec from
(select week, agent, sum((time[0]*3600+time[1]*60+time[2])) as weekly_response_time_in_sec
from(select agent_name as agent, weekofyear(date) as week, split(average_response_time,':') as
time from agent_performance) t group by agent, week)s group by agent;
```

Query ID = cloudera_20221101233535_4ba561a1-9387-4cca-a257-2a3de2da2ae0

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0027, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0027/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0027

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 23:35:17,893 Stage-1 map = 0%, reduce = 0%

2022-11-01 23:35:27,189 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.21 sec

2022-11-01 23:35:37,351 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.32 sec

MapReduce Total cumulative CPU time: 5 seconds 320 msec

Ended Job = job_1667316033891_0027

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.32 sec HDFS Read: 129398 HDFS Write: 1212 SUCCESS

Total MapReduce CPU Time Spent: 5 seconds 320 msec

OK

```
agent avg_weekly_response_time_in_sec
```

```
Abhishek 0.0
```

```
Aditya 0.0
```

```
Aditya Shinde 178.4
```

```
Aditya_iot 203.8
```

```
Amersh 0.0
```

```
Ameya Jain 126.8
```

```
Anirudh 130.6
```


Ankit Sharma 0.0

Ankitjha 26.6

Anurag Tiwari 50.6

Aravind 128.2

Ashad Nasim 231.8

Ashish 0.0

Ayushi Mishra 362.0

Bharath 160.8

Boktiar Ahmed Bappy 396.2

Chaitra K Hiremath 90.8

Deepranjan Gupta 319.2

Dibyanshu 7.6

Harikrishnan Shaji 203.8

Hitesh Choudhary 0.0

Hrisikesh Neogi 303.0

Hyder Abbas 0.0

Ineuron Intelligence 0.0

Ishawant Kumar 300.8

Jawala Prakash 565.4

Jayant Kumar 110.6

Jaydeep Dixit 266.4

Khushboo Priya 367.8

Madhulika G 398.6

Mahak 0.0

Mahesh Sarade 278.6

Maitry 383.0

Maneesh 27.0

Manjunatha A 217.0

Mithun S 173.6

Mukesh 20.0

Mukesh Rao 78.8

Muskan Garg 35.6

Nandani Gupta 359.2

Nishtha Jain 364.8

Nitin M 0.0

Prabir Kumar Satapathy 228.0

Prateek _jot 135.0

Prerna Singh 286.0

Rishav Dash	363.8
Rohan	0.0
Saif Khan	0.0
Saikumarreddy N	151.0
Samprit	0.0
Sandipan Saha	35.4
Sanjeev Kumar	307.2
Sanjeevan	0.0
Saurabh Shukla	21.0
Shiva Srivastava	60.0
Shivan K	287.4
Shivan_S	14.6
Shivananda Sonwane	336.0
Shubham Sharma	290.0
Sowmiya Sivakumar	157.0
Spuri	0.0
Sudhanshu Kumar	24.0
Suraj S Bilgi	36.4
Swati	346.8
Tarun	0.0
Uday Mishra	0.0
Vasanth P	0.0
Vivek	82.2
Wasim	178.2
Zeeshan	370.4

Time taken: 30.063 seconds, Fetched: 71 row(s)

Q13) average weekly resolution time for each agents

Solution:

```
hive> select agent, avg(weekly_resolution_time_in_sec) as avg_weekly_resolution_time_in_sec from
(select week, agent, sum((time[0]*3600+time[1]*60+time[2])) as weekly_resolution_time_in_sec
from(select agent_name as agent, weekofyear(date) as week, split(average_resolution_time,':') as
time from agent_performance) t group by agent, week)s group by agent;
```

Query ID = cloudera_20221101234040_71190205-b2a4-41ad-b17e-5be389b17f6a

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0028, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0028/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0028

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 23:40:16,339 Stage-1 map = 0%, reduce = 0%

2022-11-01 23:40:28,723 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.64 sec

2022-11-01 23:40:49,453 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 11.52 sec

MapReduce Total cumulative CPU time: 11 seconds 520 msec

Ended Job = job_1667316033891_0028

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 11.52 sec HDFS Read: 129404 HDFS Write: 1271 SUCCESS

Total MapReduce CPU Time Spent: 11 seconds 520 msec

OK

agent avg_weekly_resolution_time_in_sec

Abhishek 0.0

Aditya 0.0

Aditya Shinde 3723.8

Aditya_jot 3535.8

Amersh 0.0

Ameya Jain 1971.0

Anirudh 1112.8

Ankit Sharma 0.0

Ankitjha 326.8

Anurag Tiwari 443.6

Aravind 3213.4

Ashad Nasim 125.6

Ashish 0.0

Ayushi Mishra 5535.6

Bharath 3885.6

Boktiar Ahmed Bappy 6143.0

Chaitra K Hiremath 533.8

Deepranjan Gupta 7507.0

Dibyanshu 148.0

Harikrishnan Shaji 4070.6

Hitesh Choudhary 17.0
Hrisikesh Neogi 5557.8
Hyder Abbas 0.0
Ineuron Intelligence 0.0
Ishawant Kumar 5190.6
Jawala Prakash 4682.6
Jayant Kumar 1769.6
Jaydeep Dixit 6486.4
Khushboo Priya 6026.6
Madhulika G 5595.4
Mahak 240.8
Mahesh Sarade 3297.0
Maitry 4602.4
Maneesh 250.0
Manjunatha A 6105.2
Mithun S 2080.0
Mukesh 455.6
Mukesh Rao 2763.4
Muskan Garg 691.2
Nandani Gupta 6538.2
Nishtha Jain 3350.4
Nitin M 0.0
Prabir Kumar Satapathy 2139.8
Prateek _iot 2917.4
Perna Singh 5948.4
Rishav Dash 6114.6
Rohan 0.0
Saif Khan 0.0
Saikumarreddy N 2240.2
Samprit 20.6
Sandipan Saha 947.0
Sanjeev Kumar 6189.8
Sanjeevan 0.0
Saurabh Shukla 428.4
Shiva Srivastava 532.8
Shivan K 5709.6
Shivan_S 220.4
Shivananda Sonwane 7611.6

Shubham Sharma 6259.0

Sowmiya Sivakumar 2144.0

Spuri 0.0

Sudhanshu Kumar 701.8

Suraj S Bilgi 946.0

Swati 4419.4

Tarun 542.8

Uday Mishra 0.0

Vasanth P 0.0

Vivek 916.6

Wasim 4133.2

Zeeshan 3870.4

Time taken: 48.741 seconds, Fetched: 71 row(s)

Q14) Find the number of chats on which they have received a feedback

Solution:

hive> select agent_name as agent, sum(total_chats) as chats, sum(total_feedback) as chats_with_feedback_received from agent_performance group by agent_name;

Query ID = cloudera_20221101234747_993bf134-206c-4ae0-9001-cba22ef8a2a8

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job_1667316033891_0029, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0029/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0029

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-01 23:47:53,761 Stage-1 map = 0%, reduce = 0%

2022-11-01 23:48:06,087 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.01 sec

2022-11-01 23:48:19,713 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 6.5 sec

MapReduce Total cumulative CPU time: 6 seconds 500 msec

Ended Job = job_1667316033891_0029

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.5 sec HDFS Read: 125793 HDFS Write: 1285 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 500 msec

OK

agent chats chats_with_feedback_received

Abhishek 0 0

Aditya 0 0

Aditya Shinde 277 153

Aditya_iot 231 131

Amersh 0 0

Ameya Jain 322 228

Anirudh 81 39

Ankit Sharma 0 0

Ankitjha 5 3

Anurag Tiwari 4 3

Aravind 366 233

Ashad Nasim 18 9

Ashish 0 0

Ayushi Mishra 514 329

Bharath 369 247

Boktiar Ahmed Bappy 452 311

Chaitra K Hiremath 64 37

Deepranjan Gupta 493 312

Dibyanshu 1 0

Harikrishnan Shaji 381 231

Hitesh Choudhary 1 0

Hrisikesh Neogi 578 367

Hyder Abbas 0 0

Ineuron Intelligence 0 0

Ishawant Kumar 338 202

Jawala Prakash 439 250

Jayant Kumar 127 70

Jaydeep Dixit 512 305

Khushboo Priya 446 289

Madhulika G 469 281

Mahak 7 5

Mahesh Sarade 364 216

Maitry 542 347

Maneesh 4 3

Manjunatha A	413	254
Mithun S	503	364
Mukesh	19	17
Mukesh Rao	5	5
Muskan Garg	56	37
Nandani Gupta	560	308
Nishtha Jain	373	257
Nitin M	0	0
Prabir Kumar Satapathy	299	222
Prateek _iot	190	107
Prerna Singh	401	235
Rishav Dash	409	264
Rohan	0	0
Saif Khan	0	0
Saikumarreddy N	364	290
Samprit	1	0
Sandipan Saha	30	18
Sanjeev Kumar	507	311
Sanjeevan	0	0
Saurabh Shukla	16	8
Shiva Srivastava	53	46
Shivan K	357	243
Shivan_S	7	4
Shivananda Sonwane	441	263
Shubham Sharma	510	300
Sowmiya Sivakumar	206	141
Spuri	0	0
Sudhanshu Kumar	2	2
Suraj S Bilgi	28	15
Swati	524	302
Tarun	22	6
Uday Mishra	0	0
Vasanth P	0	0
Vivek	44	20
Wasim	433	284
Zeeshan	542	335

Time taken: 41.182 seconds, Fetched: 71 row(s)

Q15) Total contribution hour for each and every agents weekly basis

Solution:

```
hive> select week, agent, sum((time[0]*3600+time[1]*60+time[2])/3600) as total_hrs_contributed
from(select agent, weekofyear(date) as week, split(duration,':') as time from agent_logging) t group by
agent, week;
```

Query ID = cloudera_20221102000303_0195e12e-5cf9-4ba2-9b71-195b9943d4ee

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1667316033891_0032, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667316033891_0032/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667316033891_0032

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2022-11-02 00:03:43,956 Stage-1 map = 0%, reduce = 0%

2022-11-02 00:03:58,926 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.5 sec

2022-11-02 00:04:13,594 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.45 sec

MapReduce Total cumulative CPU time: 9 seconds 450 msec

Ended Job = job_1667316033891_0032

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.45 sec HDFS Read: 68662 HDFS Write: 3068 SUCCESS

Total MapReduce CPU Time Spent: 9 seconds 450 msec

OK

```
week  agent  total_hrs_contributed
```

```
30  Aditya Shinde  0.036111111111111111
```

```
29  Aditya_iot    6.095277777777778
```

```
30  Aditya_iot    9.635833333333334
```

```
30  Amersh  3.0638888888888887
```

```
29  Ameya Jain    24.083055555555553
```

```
30  Ameya Jain    17.9925
```

```
30  Ankitjha     2.2669444444444444
```

```
29  Anurag Tiwari 0.2644444444444444
```


30 Anurag Tiwari 2.514444444444445
29 Aravind 24.23555555555555
30 Aravind 0.0636111111111112
29 Ayushi Mishra 17.79027777777778
30 Ayushi Mishra 20.33138888888889
29 Bharath 24.07083333333333
30 Bharath 24.00583333333335
29 Boktiar Ahmed Bappy 17.75027777777778
30 Boktiar Ahmed Bappy 22.51833333333334
29 Chaitra K Hiremath 2.234722222222225
30 Chaitra K Hiremath 32.09083333333336
29 Deepranjan Gupta 48.99638888888889
30 Deepranjan Gupta 57.27888888888887
29 Dibyanshu 27.743888888888907
30 Dibyanshu 24.851944444444474
29 Harikrishnan Shaji 21.43833333333333
30 Harikrishnan Shaji 32.27638888888889
29 Hrisikesh Neogi 26.89138888888889
30 Hrisikesh Neogi 30.67722222222223
29 Hyder Abbas 0.3361111111111114
30 Hyder Abbas 0.0519444444444446
29 Ineuron Intelligence 1.448611111111111
29 Ishawant Kumar 25.72083333333333
30 Ishawant Kumar 26.05805555555555
29 Jawala Prakash 24.340000000000003
30 Jawala Prakash 22.06222222222225
29 Jaydeep Dixit 41.91444444444444
30 Jaydeep Dixit 17.926111111111112
29 Khushboo Priya 21.71583333333336
30 Khushboo Priya 21.84277777777778
29 Madhulika G 25.85083333333334
30 Madhulika G 20.20944444444447
29 Mahesh Sarade 25.48305555555552
30 Mahesh Sarade 17.43888888888886
29 Maitry 24.65777777777778
30 Maitry 6.28722222222222
29 Manjunatha A 18.351111111111113
30 Manjunatha A 22.92361111111111

29	Mithun S	17.3797222222222
30	Mithun S	27.7938888888889
30	Mukesh	8.905
29	Muskan Garg	3.31861111111111
30	Muskan Garg	14.0169444444444
29	Nandani Gupta	17.3338888888889
30	Nandani Gupta	22.8383333333333
29	Nishtha Jain	22.1158333333333
30	Nishtha Jain	21.7369444444444
29	Nitin M	0.79888888888889
29	Prabir Kumar Satapathy	17.5244444444444
30	Prabir Kumar Satapathy	15.8527777777778
29	Prateek _iot	7.2697222222222
30	Prateek _iot	11.1483333333333
29	Prerna Singh	18.5175000000000
30	Prerna Singh	27.1980555555556
29	Rishav Dash	18.8938888888889
30	Rishav Dash	22.8816666666667
29	Saikumarreddy N	24.9805555555556
30	Saikumarreddy N	18.1569444444444
29	Sanjeev Kumar	19.3608333333333
30	Sanjeev Kumar	25.3263888888889
29	Saurabh Shukla	16.6630555555556
29	Shiva Srivastava	1.90611111111111
30	Shiva Srivastava	13.0880555555556
29	Shivan K	16.7138888888889
30	Shivan K	19.3883333333333
29	Shivananda Sonwane	20.8341666666667
30	Shivananda Sonwane	28.4536111111111
29	Shubham Sharma	30.5102777777778
30	Shubham Sharma	23.2880555555556
29	Sowmiya Sivakumar	17.0658333333333
30	Sowmiya Sivakumar	27.6883333333333
29	Sudhanshu Kumar	24.4547222222222
30	Sudhanshu Kumar	21.7763888888889
30	Suraj S Bilgi	12.5591666666667
29	Swati	18.8586111111111
30	Swati	6.1425

```
26 Tarun 10.13888888888889
29 Wasim 19.62555555555554
30 Wasim 28.55361111111113
29 Zeeshan 24.42750000000002
30 Zeeshan 24.66111111111115

Time taken: 45.08 seconds, Fetched: 89 row(s)
```

Q16) Perform inner join, left join and right join based on the agent column and after joining the table export that data into your local system.

Solution:

Inner Join:

```
hive> select l.*, p.*
> from
> agent_logging l
> inner join
> agent_performance p
> on l.agent = p.agent_name
> limit 5;
```

Query ID = cloudera_20221102012525_d501b321-d4d3-421b-8138-c772966267be

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102012525_d501b321-d4d3-421b-8138-c772966267be.log

2022-11-02 01:25:37 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 01:25:44 Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-25-18_397_2973530280533039045-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable

2022-11-02 01:25:44 Uploaded 1 File to: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-25-18_397_2973530280533039045-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (39341 bytes)

2022-11-02 01:25:44 End of local task; Time Taken: 6.518 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0001/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0001

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 01:26:19,004 Stage-3 map = 0%, reduce = 0%

2022-11-02 01:26:37,823 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.27 sec

MapReduce Total cumulative CPU time: 4 seconds 270 msec

Ended Job = job_1667377030429_0001

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 4.27 sec HDFS Read: 13194 HDFS Write: 542 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 270 msec

OK

l.s_no	l.agent	l.date	l.login_time	l.logout_time	l.duration	p.s_no	p.date	p.agent_name	p.total_chats	p.average_response_time	p.average_resolution_time	p.average_rating	p.total_feedback
16 4.11	Perna Singh	2022-07-30	12:32:28	14:10:08	01:37:40	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	
75 4.11	Perna Singh	2022-07-29	17:47:06	21:03:44	03:16:37	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	
91 4.11	Perna Singh	2022-07-29	15:08:22	17:20:49	02:12:27	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	
110 4.11	Perna Singh	2022-07-29	12:08:23	12:11:35	00:03:11	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	
336 4.11	Perna Singh	2022-07-27	13:11:06	20:58:35	07:47:29	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	

Time taken: 81.761 seconds, Fetched: 5 row(s)

Left Join:

hive> select l.*, p.*

> from

> agent_logging l

> left join

> agent_performance p

> on l.agent = p.agent_name

> limit 5;

Query ID = cloudera_20221102012727_6495894f-c7bf-444f-a20f-ab60098d710d

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102012727_6495894f-c7bf-444f-a20f-ab60098d710d.log

2022-11-02 01:27:50 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 01:27:52 Dump the side-table for tag: 1 with group count: 71 into file: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-27-41_896_8829501887617462410-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.hashtable

2022-11-02 01:27:53 Uploaded 1 File to: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-27-41_896_8829501887617462410-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.hashtable (77631 bytes)

2022-11-02 01:27:53 End of local task; Time Taken: 3.103 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0002, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0002/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0002

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 01:28:10,612 Stage-3 map = 0%, reduce = 0%

2022-11-02 01:28:28,294 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.95 sec

MapReduce Total cumulative CPU time: 4 seconds 950 msec

Ended Job = job_1667377030429_0002

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 4.95 sec HDFS Read: 13247 HDFS Write: 600 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 950 msec

OK

l.s_no	l.agent	l.date	l.login_time	l.logout_time	l.duration	p.s_no	p.date	p.agent_name	p.total_chats	p.average_response_time	p.average_resolution_time	p.average_rating	p.total_feedback
--------	---------	--------	--------------	---------------	------------	--------	--------	--------------	---------------	-------------------------	---------------------------	------------------	------------------

1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10	69	2022-07-30	Shivananda Sonwane	4				
00:01:14	00:16:53	5.0	1										
1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10	73	2022-07-29	Shivananda Sonwane	14				
00:00:45	00:15:38	4.679											
1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10	214	2022-07-28	Shivananda Sonwane	5				
00:00:31	00:38:04	5.0	4										
1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10	285	2022-07-27	Shivananda Sonwane	26				
00:01:12	00:20:10	4.2218											
1	Shivananda Sonwane	2022-07-30	15:35:29	17:39:39	02:04:10	360	2022-07-26	Shivananda Sonwane	24				
00:00:51	00:22:28	5.0	14										

Time taken: 47.602 seconds, Fetched: 5 row(s)

Right Join:

hive> select l.*, p.*

> from

> agent_logging l

> right join

> agent_performance p

> on l.agent = p.agent_name

> limit 3;

Query ID = cloudera_20221102013131_5a7932b2-a989-4ece-8444-63b3a7a724b6

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102013131_5a7932b2-a989-4ece-8444-63b3a7a724b6.log

2022-11-02 01:31:28 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 01:31:31 Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-31-18_589_6374086923008905591-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile20--.hashtable

2022-11-02 01:31:31 Uploaded 1 File to: file:/tmp/cloudera/a7e5d834-a1dc-4221-bc58-c45c9866afe5/hive_2022-11-02_01-31-18_589_6374086923008905591-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile20--.hashtable (39341 bytes)

2022-11-02 01:31:31 End of local task; Time Taken: 2.557 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0003, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0003/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0003

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 01:31:46,028 Stage-3 map = 0%, reduce = 0%

2022-11-02 01:32:03,722 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 3.86 sec

MapReduce Total cumulative CPU time: 3 seconds 860 msec

Ended Job = job_1667377030429_0003

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 3.86 sec HDFS Read: 13221 HDFS Write: 324 SUCCESS

Total MapReduce CPU Time Spent: 3 seconds 860 msec

OK

l.s_no	l.agent	l.date	l.login_time	l.logout_time	l.duration	p.s_no	p.date	p.agent_name	p.total_chats	p.average_response_time	p.average_resolution_time	p.average_rating	p.total_feedback
16 4.11	Perna Singh	2022-07-30	12:32:28	14:10:08	01:37:40	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20	9	
75 4.11	Perna Singh	2022-07-29	17:47:06	21:03:44	03:16:37	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20		
91 4.11	Perna Singh	2022-07-29	15:08:22	17:20:49	02:12:27	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20		

Time taken: 46.342 seconds, Fetched: 3 row(s)

Export data into local system:

Inner Join:

```
[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_logging l inner join agent.agent_performance p on l.agent = p.agent_name limit 5' > /tmp/Agent_data/agent_inner_join.csv
```

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

Query ID = cloudera_20221102021313_3dbb874d-fdc2-49bd-b1d6-f3eb47b5c6d5

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102021313_3dbb874d-fdc2-49bd-b1d6-f3eb47b5c6d5.log

2022-11-02 02:13:23 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 02:13:27 Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/53d10f02-679f-4b61-918e-75477428efc2/hive_2022-11-02_02-13-11_168_5834537046314652500-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable

2022-11-02 02:13:27 Uploaded 1 File to: file:/tmp/cloudera/53d10f02-679f-4b61-918e-75477428efc2/hive_2022-11-02_02-13-11_168_5834537046314652500-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (39341 bytes)

2022-11-02 02:13:27 End of local task; Time Taken: 4.238 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0006/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0006

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 02:13:47,531 Stage-3 map = 0%, reduce = 0%

2022-11-02 02:14:05,222 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 5.09 sec

MapReduce Total cumulative CPU time: 5 seconds 90 msec

Ended Job = job_1667377030429_0006

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 5.09 sec HDFS Read: 13194 HDFS Write: 542 SUCCESS

Total MapReduce CPU Time Spent: 5 seconds 90 msec

OK

Time taken: 56.553 seconds, Fetched: 5 row(s)

Left Join:

```
[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_logging l left join agent.agent_performance p on l.agent = p.agent_name limit 5' > /tmp/Agent_data/agent_left_join.csv
```

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

Query ID = cloudera_20221102021414_f7ab8dd8-fdc3-448a-a846-512bbd2451c6

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102021414_f7ab8dd8-fdc3-448a-a846-512bbd2451c6.log

2022-11-02 02:14:54 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 02:14:57 Dump the side-table for tag: 1 with group count: 71 into file: file:/tmp/cloudera/4cf935ee-290a-4387-bccc-1d023b332b07/hive_2022-11-02_02-14-43_629_6565383743414948191-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable

2022-11-02 02:14:57 Uploaded 1 File to: file:/tmp/cloudera/4cf935ee-290a-4387-bccc-1d023b332b07/hive_2022-11-02_02-14-43_629_6565383743414948191-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable (77631 bytes)

2022-11-02 02:14:57 End of local task; Time Taken: 3.059 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0007/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0007

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 02:15:14,314 Stage-3 map = 0%, reduce = 0%

2022-11-02 02:15:31,808 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.5 sec

MapReduce Total cumulative CPU time: 4 seconds 500 msec

Ended Job = job_1667377030429_0007

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 4.5 sec HDFS Read: 13055 HDFS Write: 600 SUCCESS

Total MapReduce CPU Time Spent: 4 seconds 500 msec

OK

Time taken: 50.644 seconds, Fetched: 5 row(s)

Right Join:

```
[cloudera@quickstart ~]$ hive -e 'select l.*, p.* from agent.agent_logging l right join  
agent.agent_performance p on l.agent = p.agent_name limit 5' >  
/tmp/Agent_data/agent_right_join.csv
```

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

Query ID = cloudera_20221102021616_6c0ea584-01f3-4d15-b707-8b2116b0307d

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20221102021616_6c0ea584-01f3-4d15-b707-8b2116b0307d.log

2022-11-02 02:16:20 Starting to launch local task to process map join; maximum memory = 932184064

2022-11-02 02:16:22 Dump the side-table for tag: 0 with group count: 49 into file: file:/tmp/cloudera/7ef817ad-963c-4d3b-b113-e87e225620e4/hive_2022-11-02_02-16-08_907_5363628296010565902-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00-.hashtable

2022-11-02 02:16:23 Uploaded 1 File to: file:/tmp/cloudera/7ef817ad-963c-4d3b-b113-e87e225620e4/hive_2022-11-02_02-16-08_907_5363628296010565902-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00-.hashtable (39341 bytes)

2022-11-02 02:16:23 End of local task; Time Taken: 3.068 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0008/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0008

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2022-11-02 02:16:44,458 Stage-3 map = 0%, reduce = 0%

2022-11-02 02:16:59,648 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.45 sec

16 4.11	Perna Singh 9	2022-07-30	12:32:28	14:10:08	01:37:40	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20
75 4.11	Perna Singh 9	2022-07-29	17:47:06	21:03:44	03:16:37	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20
91 4.11	Perna Singh 9	2022-07-29	15:08:22	17:20:49	02:12:27	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20
110 4.11	Perna Singh 9	2022-07-29	12:08:23	12:11:35	00:03:11	1	2022-07-30	Perna Singh	11	00:00:38	00:04:20

Q17) Perform partitioning on top of the agent column and then on top of that perform bucketing for each partitioning.

Solution:

First set the below mentioned properties to be true.

```
hive> set hive.exec.dynamic.partition=true;
```

```
hive> set hive.exec.dynamic.patition.mode=nonstrict;
```

Create Partition_bucketed table:

```
hive> create table partition_bucketed_logging
```

```
> (  
> s_no int,  
> date date,  
> login_time string,  
> logout_time string,  
> duration string  
> )  
> partitioned by (agent string)  
> clustered by(s_no)  
> into 4 buckets  
> row format delimited  
> fields terminated by ','  
> stored as textfile;
```

OK

Time taken: 0.156 seconds

Load data into Patition_bucketed table:

```
hive> insert overwrite table partition_bucketed_logging partition(agent) select s_no, date, login_time,  
logout_time, duration, agent from agent_logging;
```

Query ID = cloudera_20221102040808_6531ce58-de2e-4c4a-a582-331557eb8787

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1667377030429_0010, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667377030429_0010/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667377030429_0010

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2022-11-02 04:08:54,135 Stage-1 map = 0%, reduce = 0%

2022-11-02 04:09:15,365 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.32 sec

MapReduce Total cumulative CPU time: 6 seconds 320 msec

Ended Job = job_1667377030429_0010

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/agent.db/partition_bucketed_logging/.hive-staging_hive_2022-11-02_04-08-38_814_1998316040267771961-1/-ext-10000

Loading data to table agent.partition_bucketed_logging partition (agent=null)

Time taken for load dynamic partitions : 11817

Loading partition {agent=Sudhanshu Kumar}

Loading partition {agent=Aditya Shinde}

Loading partition {agent=Suraj S Bilgi}

Loading partition {agent=Hrisikesh Neogi}

Loading partition {agent=Dibyanshu}

Loading partition {agent=Shiva Srivastava}

Loading partition {agent=Saikumarreddy N}

Loading partition {agent=Jaydeep Dixit}

Loading partition {agent=Ankitjha}

Loading partition {agent=Prabir Kumar Satapathy}

Loading partition {agent=Deepranjan Gupta}

Loading partition {agent=Khushboo Priya}

Loading partition {agent=Hyder Abbas}

Loading partition {agent=Amersh}

Loading partition {agent=Anurag Tiwari}

Loading partition {agent=Madhulika G}

Loading partition {agent=Aravind}

Loading partition {agent=Prateek _iot}

Loading partition {agent=Mithun S}

Loading partition {agent=Mahesh Sarade}

Loading partition {agent=Ameya Jain}

Loading partition {agent=Swati}

Loading partition {agent=Bharath}

Loading partition {agent=Sowmiya Sivakumar}

Loading partition {agent=Ineuron Intelligence}

Loading partition {agent=Prerna Singh}
Loading partition {agent=Wasim}
Loading partition {agent=Zeeshan}
Loading partition {agent=Manjunatha A}
Loading partition {agent=Nishtha Jain}
Loading partition {agent=Harikrishnan Shaji}
Loading partition {agent=Shubham Sharma}
Loading partition {agent=Boktiar Ahmed Bappy}
Loading partition {agent=Saurabh Shukla}
Loading partition {agent=Rishav Dash}
Loading partition {agent=Jawala Prakash}
Loading partition {agent=Mukesh}
Loading partition {agent=Tarun}
Loading partition {agent=Sanjeev Kumar}
Loading partition {agent=Maitry}
Loading partition {agent=Nandani Gupta}
Loading partition {agent=Aditya_iot}
Loading partition {agent=Chaitra K Hiremath}
Loading partition {agent=Shivan K}
Loading partition {agent=Shivananda Sonwane}
Loading partition {agent=Nitin M}
Loading partition {agent=Ishawant Kumar}
Loading partition {agent=Muskan Garg}
Loading partition {agent=Ayushi Mishra}

Time taken for adding to write entity : 31

Partition agent.partition_bucketed_logging{agent=Aditya Shinde} stats: [numFiles=1, numRows=1, totalSize=42, rawDataSize=41]
Partition agent.partition_bucketed_logging{agent=Aditya_iot} stats: [numFiles=1, numRows=9, totalSize=377, rawDataSize=368]
Partition agent.partition_bucketed_logging{agent=Amersh} stats: [numFiles=1, numRows=4, totalSize=168, rawDataSize=164]
Partition agent.partition_bucketed_logging{agent=Ameya Jain} stats: [numFiles=1, numRows=10, totalSize=420, rawDataSize=410]
Partition agent.partition_bucketed_logging{agent=Ankitjha} stats: [numFiles=1, numRows=4, totalSize=168, rawDataSize=164]
Partition agent.partition_bucketed_logging{agent=Anurag Tiwari} stats: [numFiles=1, numRows=37, totalSize=1553, rawDataSize=1516]
Partition agent.partition_bucketed_logging{agent=Aravind} stats: [numFiles=1, numRows=10, totalSize=420, rawDataSize=410]
Partition agent.partition_bucketed_logging{agent=Ayushi Mishra} stats: [numFiles=1, numRows=18, totalSize=755, rawDataSize=737]
Partition agent.partition_bucketed_logging{agent=Bharath} stats: [numFiles=1, numRows=9, totalSize=378, rawDataSize=369]
Partition agent.partition_bucketed_logging{agent=Boktiar Ahmed Bappy} stats: [numFiles=1, numRows=17, totalSize=709, rawDataSize=692]
Partition agent.partition_bucketed_logging{agent=Chaitra K Hiremath} stats: [numFiles=1, numRows=13, totalSize=543, rawDataSize=530]
Partition agent.partition_bucketed_logging{agent=Deepranjan Gupta} stats: [numFiles=1, numRows=58, totalSize=2433, rawDataSize=2375]
Partition agent.partition_bucketed_logging{agent=Dibyanshu} stats: [numFiles=1, numRows=208, totalSize=8719, rawDataSize=8511]

Partition agent.partition_bucketed_logging{agent=Harikrishnan Shaji} stats: [numFiles=1, numRows=23, totalSize=963, rawDataSize=940]

Partition agent.partition_bucketed_logging{agent=Hrisikesh Neogi} stats: [numFiles=1, numRows=37, totalSize=1544, rawDataSize=1507]

Partition agent.partition_bucketed_logging{agent=Hyder Abbas} stats: [numFiles=1, numRows=2, totalSize=84, rawDataSize=82]

Partition agent.partition_bucketed_logging{agent=Ineuron Intelligence} stats: [numFiles=1, numRows=1, totalSize=42, rawDataSize=41]

Partition agent.partition_bucketed_logging{agent=Ishawant Kumar} stats: [numFiles=1, numRows=49, totalSize=2052, rawDataSize=2003]

Partition agent.partition_bucketed_logging{agent=Jawala Prakash} stats: [numFiles=1, numRows=16, totalSize=668, rawDataSize=652]

Partition agent.partition_bucketed_logging{agent=Jaydeep Dixit} stats: [numFiles=1, numRows=11, totalSize=459, rawDataSize=448]

Partition agent.partition_bucketed_logging{agent=Khushboo Priya} stats: [numFiles=1, numRows=18, totalSize=752, rawDataSize=734]

Partition agent.partition_bucketed_logging{agent=Madhulika G} stats: [numFiles=1, numRows=17, totalSize=713, rawDataSize=696]

Partition agent.partition_bucketed_logging{agent=Mahesh Sarade} stats: [numFiles=1, numRows=36, totalSize=1509, rawDataSize=1473]

Partition agent.partition_bucketed_logging{agent=Maitry} stats: [numFiles=1, numRows=5, totalSize=210, rawDataSize=205]

Partition agent.partition_bucketed_logging{agent=Manjunatha A} stats: [numFiles=1, numRows=8, totalSize=333, rawDataSize=325]

Partition agent.partition_bucketed_logging{agent=Mithun S} stats: [numFiles=1, numRows=14, totalSize=586, rawDataSize=572]

Partition agent.partition_bucketed_logging{agent=Mukesh} stats: [numFiles=1, numRows=3, totalSize=124, rawDataSize=121]

Partition agent.partition_bucketed_logging{agent=Muskan Garg} stats: [numFiles=1, numRows=12, totalSize=503, rawDataSize=491]

Partition agent.partition_bucketed_logging{agent=Nandani Gupta} stats: [numFiles=1, numRows=11, totalSize=458, rawDataSize=447]

Partition agent.partition_bucketed_logging{agent=Nishtha Jain} stats: [numFiles=1, numRows=18, totalSize=754, rawDataSize=736]

Partition agent.partition_bucketed_logging{agent=Nitin M} stats: [numFiles=1, numRows=1, totalSize=42, rawDataSize=41]

Partition agent.partition_bucketed_logging{agent=Prabir Kumar Satapathy} stats: [numFiles=1, numRows=26, totalSize=1091, rawDataSize=1065]

Partition agent.partition_bucketed_logging{agent=Prateek _iot} stats: [numFiles=1, numRows=17, totalSize=711, rawDataSize=694]

Partition agent.partition_bucketed_logging{agent=Prerna Singh} stats: [numFiles=1, numRows=18, totalSize=753, rawDataSize=735]

Partition agent.partition_bucketed_logging{agent=Rishav Dash} stats: [numFiles=1, numRows=12, totalSize=504, rawDataSize=492]

Partition agent.partition_bucketed_logging{agent=Saikumarreddy N} stats: [numFiles=1, numRows=10, totalSize=420, rawDataSize=410]

Partition agent.partition_bucketed_logging{agent=Sanjeev Kumar} stats: [numFiles=1, numRows=20, totalSize=839, rawDataSize=819]

Partition agent.partition_bucketed_logging{agent=Saurabh Shukla} stats: [numFiles=1, numRows=40, totalSize=1680, rawDataSize=1640]

Partition agent.partition_bucketed_logging{agent=Shiva Srivastava} stats: [numFiles=1, numRows=15, totalSize=629, rawDataSize=614]

Partition agent.partition_bucketed_logging{agent=Shivan K} stats: [numFiles=1, numRows=36, totalSize=1506, rawDataSize=1470]

Partition agent.partition_bucketed_logging{agent=Shivananda Sonwane} stats: [numFiles=1, numRows=15, totalSize=625, rawDataSize=610]

Partition agent.partition_bucketed_logging{agent=Shubham Sharma} stats: [numFiles=1, numRows=35, totalSize=1469, rawDataSize=1434]

Partition agent.partition_bucketed_logging{agent=Sowmiya Sivakumar} stats: [numFiles=1, numRows=24, totalSize=1005, rawDataSize=981]

Partition agent.partition_bucketed_logging{agent=Sudhanshu Kumar} stats: [numFiles=1, numRows=11, totalSize=462, rawDataSize=451]

Partition agent.partition_bucketed_logging{agent=Suraj S Bilgi} stats: [numFiles=1, numRows=5, totalSize=206, rawDataSize=201]

Partition agent.partition_bucketed_logging{agent=Swati} stats: [numFiles=1, numRows=5, totalSize=210, rawDataSize=205]

Partition agent.partition_bucketed_logging{agent=Tarun} stats: [numFiles=1, numRows=1, totalSize=43, rawDataSize=42]

Partition agent.partition_bucketed_logging{agent=Wasim} stats: [numFiles=1, numRows=20, totalSize=840, rawDataSize=820]

Partition agent.partition_bucketed_logging{agent=Zeeshan} stats: [numFiles=1, numRows=10, totalSize=419, rawDataSize=409]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 6.32 sec HDFS Read: 61313 HDFS Write: 45367 SUCCESS

Total MapReduce CPU Time Spent: 6 seconds 320 msec

OK

Time taken: 57.42 seconds